

EB-JEPA for EEG — What We Did and What It Shows

A precise but compact summary of our original contributions

EB-JEPA EEG track

June 2026

Setting. All experiments use TUAB v3.0.1 (2717 train / 276 eval recordings, 19 channels, 200 Hz, 10-second windows, patient-disjoint split). The encoder is pretrained with SSL (no labels), then a frozen logistic regression probe is trained on the labelled set. Primary metric: per-recording BACC (mean-pool 16 windows per patient).

1 What the framework provides (not our contribution)

The **EB-JEPA library** provides the two-view SSL training loop, the VICReg and SIGReg loss functions, the Projector MLP, the RNNPredictor, and the data loading pipeline.

The **baseline Conv1D encoder** (4 blocks of strided Conv1d, BatchNorm, GELU, 0.4M parameters) produces a per-window embedding via global average pooling of the final feature map.

Standard augmentations (Gaussian noise $\sigma=0.1$, amplitude jitter $\pm 20\%$, channel drop $p=0.2$) are the default pipeline.

The baseline with these defaults and VICReg scores **0.796 BACC per-recording**.

2 Contribution 1 — Corruption augmentation (+2.3 pp, our main lever)

What we added. On top of the standard augmentations, each view receives:

- 20 % of time steps **masked to zero** (contiguous or random)
- 20 % of time steps replaced by $\pm 6\sigma$ **spike outliers**

This is inspired by the observation that real EEG artefacts (electrode pop, muscle noise, electrode drift) produce exactly these patterns. The encoder is forced to learn what persists through corruption — the clinically meaningful signal structure — rather than low-level amplitude statistics.

What it shows.

Augmentation	BACC _{REC} (3-seed)	AUROC
Standard only	0.796	0.888
+ Corruption	0.819 $\pm$.004	0.900 $\pm$.006

The gain is $+0.023$, reproducible across 3 seeds (± 0.004). No other modification — more data, wider encoder, longer training, different objective — exceeds this. Corruption is the single most impactful design decision in our work.

3 Contribution 2 — Novel finding: EEG prefers weak VICReg invariance

What VICReg invariance means. The VICReg loss has three terms:

$$\mathcal{L} = \lambda_{\text{inv}} \|z_1 - z_2\|^2 + \lambda_{\text{var}} V(z_1, z_2) + \lambda_{\text{cov}} C(z_1, z_2)$$

The invariance term $\|z_1 - z_2\|^2$ pushes the two views together. The variance term V prevents collapse by keeping each output dimension spread out. The covariance term C decorrelates dimensions. The original VICReg paper (Bardes et al., ICLR 2022, vision SSL) uses $\lambda_{\text{inv}} = 25$. The EB-JEPA library default is $\lambda_{\text{inv}} = 1$.

What we found. Running the same setup (corruption, Conv1D) with both settings, 3 seeds each:

λ_{inv}	BACC _{REC}	AUROC
1 (default)	0.819 $\pm$.004	0.900 $\pm$.006
5	0.820	0.893
10	0.814	0.900
25 (paper-correct)	0.789 $\pm$.005	0.882 $\pm$.004

Endpoints 3-seed; intermediates single-seed.

The sweet spot is $\lambda \in \{1, 5\}$ (0.819–0.820, essentially tied), with a clear drop at $\lambda=10$ and a large drop at $\lambda=25$ (**3 pp**, $\approx 6\sigma$). The ordering is not strictly monotone at the top but the overall trend is unambiguous: more invariance $\Rightarrow$ lower BACC for EEG.

Why this matters. With $\lambda_{\text{inv}} = 1$, the anti-collapse terms (variance + covariance) dominate the invariance term. The encoder is pushed to *spread representations* more than it is pushed to make two views identical. For short-window clinical EEG with a small conv encoder, this regime produces more discriminative features.

This is not a bug fix — it is a finding. To our knowledge, no prior work has systematically ablated VICReg’s invariance coefficient for EEG SSL.

4 Contribution 3 — NaiveJEPA: collapse-free prediction without regularisation

What we built. NaiveJEPA splits a single EEG window in time at ratio 0.5:

- Context ($t = 0 \rightarrow T/2$): online encoder $\rightarrow$ mean pool $\rightarrow$ MLP predictor
- Target ($t = T/2 \rightarrow T$): EMA copy of encoder $\rightarrow$ mean pool (no gradient)
- Loss: $\text{MSE}(\hat{z}_{\text{tgt}}, z_{\text{tgt}})$

Collapse is made very unlikely, but the reason is not the temporal split alone. The common claim is: “if the encoder collapses to a constant c , the MSE loss becomes the variance of the target distribution (non-zero), so collapse is prevented.” This is *wrong*: both the online and EMA encoders share the same underlying function; if that function collapses to c , both

outputs are c , the predictor can learn $\text{MLP}(c) = c$, and $\text{MSE}(c, c) = 0$ — collapse is a valid global minimum.

What actually prevents collapse is a combination of three factors:

1. **BatchNorm (dominant).** The Conv1D encoder has `BatchNorm1d` layers. If all outputs in a batch are identical, the batch variance is zero and BatchNorm divides by $\approx \varepsilon$, causing numerical explosion. True constant collapse is practically impossible with BatchNorm.
2. **EMA lag (structural asymmetry).** The target encoder is a slowly-moving copy of the online encoder (momentum ≈ 0.996). At every step, target $\neq$ online, so the predictor always has a non-stationary, non-collapsed target to chase. This is the BYOL argument: asymmetry prevents the stable collapsed equilibrium.
3. **Temporal disjointness (helpful, not sufficient).** At the start of training, first-half and second-half representations differ in content, so the loss provides useful gradient. But if both encoders eventually converge to the same constant function, this advantage disappears.

The net result: **no explicit VICReg/SIGReg regularisation is needed in practice**, and NaiveJEPA trains stably. But the architectural guarantee is provided by BatchNorm + EMA, not by the temporal split alone.

What it shows.

Objective	BACC _{REC}	Hyperparameters
VICReg + corruption	$0.819 \pm .004$	3 loss weights + EMA + augmentation policy
NaiveJEPA	0.779	EMA only (1 hyperparameter)

The gap to VICReg is only 0.040 pp with a fraction of the design complexity. NaiveJEPA also achieves **0.389 BACC on TUEV** (6-class transfer), beating ST-Transformer (0.3385) and CNN-Transformer (0.4100) trained end-to-end on TUEV — without ever seeing a TUEV label.

5 Contribution 4 — Cross-task transfer: TUEV without TUEV labels

What we did. After SSL pretraining on TUAB (binary, no labels used), we froze the encoder and trained a logistic regression probe on TUEV (6-class event detection): spike-and-slow-wave, generalised periodic discharge, lateralised periodic discharge, eye movement, artefact, background.

What it shows.

Method	TUEV BACC	Trained on TUEV?
ST-Transformer (supervised)	0.3385	Yes, end-to-end
CNN-Transformer (supervised)	0.4100	Yes, end-to-end
VICReg (TUAB SSL + probe)	0.382	Probe only
NaiveJEPA (TUAB SSL + probe)	0.389	Probe only
FFCL (supervised)	0.4551	Yes, end-to-end
BIOT pretrained (supervised)	0.5009	Yes, end-to-end

Our SSL encoder beats two fully-supervised models on the target task *while never having trained on the target task at all*. This is the strongest evidence that the representations are general: they encode structure of EEG that is relevant beyond the binary abnormality task.

6 Failure modes we documented

EEGPT architecture conflict

EEGPT’s cross-attention pools 19 channels into 1 token per time step, leaving only 10 temporal tokens. Channel-drop augmentation randomly zeroes out channels in view v_1 but not v_2 , making the two views structurally incompatible *before the transformer sees them*. The encoder cannot learn invariance to a transformation that changes the token count. Result: representations below random initialisation. Fix: remove channel-drop $\rightarrow$ TUEV BACC 0.182 $\rightarrow$ 0.358.

Evaluation protocol mismatch

Per-recording BACC is ≈ 5 pp higher than per-window for the same encoder. Most published literature (BIOT, LaBraM, FEMBA) reports per-window. We initially compared our per-recording 0.836 to LaBraM’s per-window 0.814 and concluded we matched LaBraM — this is invalid. Our best per-window score is 0.775 (SIGReg + corruption), 3.9 pp below LaBraM-Base.

Spectral loss confound

An auxiliary DDSF-style multi-scale spectral loss appeared to give +0.040 BACC when added to VICReg (0.796 $\rightarrow$ 0.836). A sweep of the spectral coefficient $\{0.05, 0.10, 0.20, 0.30\}$ shows the gain disappears: all coefficients give 0.795–0.813, within single-seed noise of 0.819 (no spectral). The original 0.836 was confounded: single seed, and the coefficient was chosen by looking at the evaluation set.

7 What the binding constraint is

Every Conv1D variant plateaus at per-recording BACC ≈ 0.82 , per-window ≈ 0.77 , regardless of:

- Encoder depth or width (scaling)
- $4\times$ more pretraining data (multi-corpus)
- Longer pretraining (not tested beyond 20 epochs, but no convergence gap)
- SSL objective (VICReg, SIGReg, NaiveJEPA all plateau near the same level)

The encoder itself is the ceiling. The 3.9 pp gap to LaBraM-Base is not a loss function problem or a data volume problem — LaBraM uses a learned VQ-codebook patch tokeniser pretrained on 2500 hours of EEG across 12 institutions. That tokeniser quality, not any SSL objective, is what drives the gap.